

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/224257163>

Improving CUDASW++ , a parallelization of Smith-Waterman for CUDA enabled devices

Conference Paper · June 2011
DOI: 10.1109/IPDPS.2011.182 · Source: IEEE Xplore

CITATIONS
27

READS
100

5 authors, including:



Doug Hains
Colorado State University
8 PUBLICATIONS 128 CITATIONS

SEE PROFILE



A. P. Wim Böhm
Colorado State University
109 PUBLICATIONS 1,368 CITATIONS

SEE PROFILE



Sanjay V. Rajopadhye
Colorado State University
167 PUBLICATIONS 2,080 CITATIONS

SEE PROFILE

Some of the authors of this publication are also working on these related projects:



Analytical Cost Models - Days of Future Past [View project](#)



Coarse-grained Hybrid Dataflow [View project](#)

Improving CUDASW++, a Parallelization of Smith-Waterman for CUDA Enabled Devices

Doug Hains, Zach Cashero, Mark Ottenberg,
Wim Bohm and Sanjay Rajopadhye

April 28, 2011

Database Sequence Alignment

- ▶ One or more *Query* sequences
- ▶ Compared to a database of sequences
- ▶ Want to know only the maximum similarity scores
- ▶ Performance measured in Giga-cell updates per seconds (GCUPS)
- ▶ CUDASW++ uses Smith-Waterman using affine gap penalties
 - ▶ ρ gap open penalty
 - ▶ σ gap extension penalty
- ▶ We identified a bottleneck in CUDASW++ and developed a wavefront tiling strategy.
- ▶ Reduces the effect of the bottleneck and increases the overall performance of the algorithm.

Smith-Waterman


- Using an affine gap penalty we have the following dependencies:

$$E_{i,j} = \max \left\{ \begin{array}{l} E_{i,j-1} - \sigma \\ H_{i,j-1} - \rho \end{array} \right\}$$

$$F_{i,j} = \max \left\{ \begin{array}{l} F_{i-1,j} - \sigma \\ H_{i-1,j} - \rho \end{array} \right\}$$

$$H_{i,j} = \max \left\{ \begin{array}{l} 0 \\ E_{i,j} \\ F_{i,j} \\ H_{i-1,j-1} + w(q_i, d_j) \end{array} \right\}$$


- $H_{i,0} = H_{0,j} = E_{i,0} = F_{0,j} = 0$



Overview of CUDASW++

- ▶ Two different tasks (kernels): Inter-task and Intra-task
 - ▶ Tasks chosen on basis of database sequence length
 - ▶ Default threshold of 3072
-
- | | |
|---|---|
| <ul style="list-style-type: none">▶ Inter-task<ul style="list-style-type: none">▶ Database seq. length < 3072▶ One thread for each sequence alignment▶ Optimizations to reduce global memory traffic▶ Dependencies stored locally | <ul style="list-style-type: none">▶ Intra-task<ul style="list-style-type: none">▶ Database seq. length ≥ 3072▶ One thread block for each sequence alignment▶ High amount of global memory traffic▶ Dependencies stored globally |
|---|---|

Intra-task is a Bottleneck




- ▶ Ran CUDASW++ on UniProtKB/Swiss-Prot database (500,000+ sequences), varied the threshold so more sequences are run with Intra-task
- ▶ By increasing the use of intra-task slightly, the performance goes down considerably
- ▶ So, why is intra-task needed?

Inter-task memory structure: Motivation for intra-task


- ▶ First, determine the number of database sequences compared to a query sequence per kernel call. (Based on device)
- ▶ Sequences are sorted by length and packed into groups in global memory.
- ▶ Size of the group is the total number of threads to be launched by the inter-task kernel. (1:1 thread to database sequence ratio)
- ▶ Sequentially launch kernels until all groups are compared.
- ▶ **Cannot make the next kernel call until all threads from previous kernel call are done.**

Load balancing in Inter-task

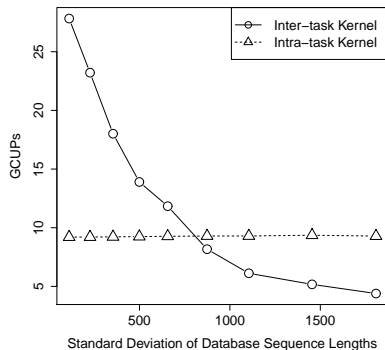
- ▶ Best Case:



- ▶ Worst Case:




Intra-task is not Affected




- ▶ Generated multiple databases of random sequences.
- ▶ Used a log-normal distribution to shift distribution of sequence lengths from best-case to worst-case.
- ▶ Intra-task is faster when distribution of lengths is uneven
- ▶ Can we improve it?

Original Intra-task



- ▶ Table is computed by diagonals which are stored in Global memory.
- ▶ Requires a number of global read/writes proportional to the length of the diagonal.
- ▶ Idle threads when diagonal length is not a multiple of thread block size.

Improving Intra-Task



- Use tiling to:
 - Reduce global memory read/writes
 - Reduce idle threads (Threads only idle during pipeline fill/empty)

Initial Results

Used a randomly generated database with all sequence lengths the threshold.

- ▶ Original CUDASW++ : 1.5 GcupS

Initial Results

Used a randomly generated database with all sequence lengths the threshold.

- ▶ Original CUDASW++ : 1.5 GcupS
- ▶ “Improved” CUDASW++ : 1.5 GcupS

Local Memory Usage - Shallow Copy

- ▶ Original CUDASW++: 0 bytes lmem
- ▶ Improved CUDASW++: 64 bytes lmem

```
int *tmp;
int h1[4], h2[4];
while(...) { /* Loop over tiles */
    ...
    tmp = h1;
    h1 = h2;
    h2 = tmp;
}
```

Local Memory Usage - Texture Memory

```
#define TILE_WIDTH 4
int tmpH[TILE_WIDTH];


for(k = 0; k < TILE_WIDTH; k++) { /* Loop over columns */
    db_char = tex2d(...); /* read db symbol */
    ...
    tmpH[k] = ...
}
```

- ▶ Loops containing texture memory calls cannot be unrolled

Use the Query Profile


- ▶ Query Profile (from Rognes, also used in Inter-task)
 - ▶ Similarity scores (w) are precomputed, storing 4 scores in 4 bytes.
- ▶ Before Query Profile
 - ▶ Each query symbol read from global memory.
 - ▶ Similarity scores read from lookup table in shared memory.
- ▶ After Query Profile
 - ▶ Retrieves four similarity scores directly per global read.
 - ▶ No shared memory reads required.

Thread Block Size and Tile Height





- ▶ Strip height: $\text{tile height} \times \text{thread block size}$
- ▶ Here, $\text{strip height} = 4 \times 3 = 12$
- ▶ If query sequence $>$ strip height, intermediate results stored in global memory.
- ▶ Found strip height of 512 to be optimal on C1060 and 1024 on C2050.
- ▶ Variations on thread block size/tile height with equal strip heights had similar performance.

Tile Width



- ▶ Tile Width does not change the number of global/shared memory accesses.
- ▶ Tile Width does increase the pipeline latency.
- ▶ Tile Width of one is optimal.

Results




Kernel	Global Memory Transactions	
	Query Len. 567	Query Len. 5478
Imp. Kernel	13,828	4,233,197
Orig. Kernel	28,345,473	291,179,739


Table: Number of total global memory accesses performed by both kernels on queries of two different sizes against the Swissprot database.

Is Global Memory Caching (Fermi) Comparable?

► Cache turned on for C2050



► Cache turned off for C2050



- Although we show an increase in performance of the original CUDASW++ due to caching, it is not as good as explicitly using shared memory.

Conclusions and Future Work

► Conclusions

- Intra-task was a bottleneck in CUDASW++
- It is necessary due to load balancing issues with inter-task
- Our wavefront tiling method improves performance on both C2050 and C1060 GPUs

► Future work

- Integrate our improvements with official version of CUDASW++
- Explore other ideas for performance tuning